An Alignment Check for Debiasing Corrections: A Negative-Control Case Study in Drug Repurposing

Rishik Kondadadi[1] and John Ortega[2]

[1] University of Minnesota
[2] Northeastern University

Abstract. A bias-diagnostic metric can judge a debiasing correction only if it is sensitive to what that correction changes; two corrections with different mechanisms can target different quantities, and a metric tuned to one will misjudge the other regardless of which is substantively better. We formalize this as an alignment check: score each correction's own penalty term as a standalone classifier of whatever the bias metric flags, and treat near-chance AUROC as evidence the metric cannot see that mechanism at all. We develop the check through a case study in drug repurposing, where a model trained to rank candidates for one indication also learns what is shared across drugs on the same organ system; in our data the convulsant bemegride ranks first among non-antiepileptic candidates for epilepsy. We test two corrections for this confound, naive area-score subtraction and a negative-control method adapted from genomics and pharmacoepidemiology, across depression, antipsychotic, and epilepsy rankings. The naive correction wins on every bias-diagnostic metric we measured, which looks like a negative result for the principled method until the alignment check is applied: its penalty term carries essentially no information (0.50 AUROC) about the compounds our bias metrics flag, while naive subtraction's term separates them well above chance (0.60–0.79). A shrinkage follow-up, a calibration check, and a direct test of the method's own identifiability assumption rule out noise and misspecification as explanations. We then close the loop formally: the correction's penalty term is, by construction, exactly orthogonal to the representation its base model is fit on, and we confirm empirically that the bias metrics' flags are themselves predictable from that same representation (0.51–0.72 AUROC). Together these explain the result without appeal to this dataset's specifics, and the same two-step check, orthogonality by construction plus empirical flag-predictability, applies to any comparison between a residualized correction and a membership-based bias metric.

1 Introduction

A bias-diagnostic metric is itself a measurement, and a measurement has a detection limit. When two debiasing corrections are compared against one bias metric and one wins, the usual reading is that the winning correction removes the bias

better. A second reading is available whenever the two corrections work through different mechanisms: the metric may simply be built to see one mechanism and not the other, in which case the comparison never tested what it looked like it was testing. Telling these readings apart requires checking, directly, whether the metric could have detected the losing correction's effect at all, rather than inferring quality from the scoreboard alone. We call this an alignment check: score each correction's own penalty term, the quantity it subtracts or projects out to produce its debiased score, as a standalone classifier of whatever the metric flags. A term that classifies the flag near chance is one the metric cannot reward or penalize, independent of whether the correction behind it is good.

We develop this check through a case study chosen because it produces exactly the puzzle the check is built to resolve. Chemical structure is not the only thing a repurposing model can learn to exploit. A model trained to separate antiepileptic drugs from everything else also, incidentally, learns what distinguishes neurology drugs from the rest of the pharmacopeia, and that second signal rides along with the first in every prediction. The practical consequence shows up at the top of a ranked candidate list: in a plain ECFP4 fingerprint model trained on known antiepileptic drugs, bemegride, a convulsant once used to induce seizures for electroconvulsive therapy, ranks first among 1,522 non-antiepileptic candidates, ahead of every genuinely plausible repurposing lead. We reproduce this finding exactly (score 0.7942, matching a prior run) as this paper's motivating case. A companion measurement study from this group audited evaluation-protocol choices in LINCS-based repurposing models and reported this same ranking [12]; that study proposes no model and leaves correction methods to future work.

The same pattern appears in aggregate, not in one example alone. Restricting to drugs with a WHO ATC classification, candidates flagged as nervous-system drugs make up 18.8% of the top 50 false positives for a depression model and 15.2% for an antipsychotic model, against base rates of 10.2% and 10.6%. The enrichment, roughly 1.5 to 1.8 times base rate, is smaller than a name-matched Hub label had suggested; ATC covers only 9.7% of the LINCS library, mostly marketed drugs, a narrower, more defensible population, used throughout.

Two corrections suggest themselves for this confound. The simple one subtracts a generic area score from the disease-specific score, directly penalizing confounded-class membership; prior informal analysis assumed this collapses precision too far to use. The more principled one borrows an idea from genomics and pharmacoepidemiology: estimate the shared organ-system factor from negative-control indications, diseases in the same area treated by a different mechanism, and remove only that factor rather than a blunt membership score. We built and evaluated the principled version. It did not beat the simple one on any bias-diagnostic metric we measured, which is where this paper could have stopped and reported a negative result.

Instead we apply the alignment check, and it changes the conclusion. Section 4 shows this orthogonality is exact only without regularization; we measure, not assume, how close the fits used here come. Section 5.7 shows the bias met-

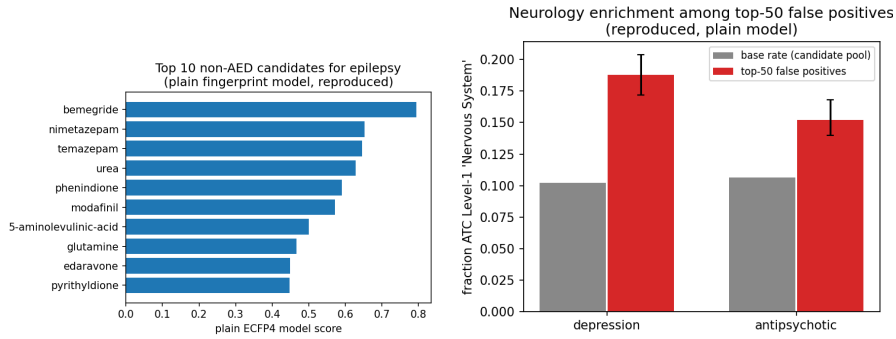


Fig. 1. Left: top 10 non-antiepileptic candidates ranked by a plain ECFP4 model trained on known antiepileptic drugs; bemegride, a convulsant, ranks first. Right: nervous-system drugs (ATC Level-1 'N') are enriched among the top-50 false positives of depression and antipsychotic models relative to each pool's base rate.

rics' flags are nonetheless predictable from the same base representation, while naive's penalty term carries no orthogonality constraint at all, which is why it tracks the flags well. Table 2's comparison is real, but it is evidence this metric cannot see the variation the principled method removes, not that the method fails to remove it.

2 Related Work

Negative controls for unwanted variation. Removing unwanted variation (RUV) via negative controls is an established genomics technique: RUV-2 [8] estimates latent confounding factors from known-null control genes and removes them before testing; later work [9] unifies several RUV variants in one factor-analysis framework. Surrogate variable analysis [16] and empirical-Bayes batch correction [11] address related but distinct versions of the problem; ComBat needs known batch labels rather than negative controls, why we follow the negative-control strategy instead. Negative-control exposures and outcomes in pharmacoepidemiology [18] formalize the same logic for observational drug-effect confounding, in active use [23]. Our area-factor step is a direct structural analogue, applied to indications rather than genes or exposures; only the application to cross-indication repurposing confounding is new.

Debiasing drug-repurposing and DTI models. NRBdMF [1] factorizes a signed drug×disease matrix from SIDER [14] (-1 indications, $+1$ side effects) to reduce false positives through within-drug directionality; it never uses other diseases as negative controls, so we include it as a baseline. TAPB [17] and DRDM [19] address other biases (target-popularity, graph-degree) in drug-target/repositioning prediction; neither targets organ-system confounding or uses negative controls. SIDER-derived phenotypic similarity was first shown to carry repositioning-relevant signal by Campillos et al. [3,13], the basis for our hard negatives.

Benchmarks and resources. repoDB [2] provides a curated gold-standard set of approved and failed drug-disease pairs and is a standard secondary comparison point in this literature. The Drug Repurposing Hub [5], source of a name-matched therapeutic-area label used in prior work from this group, has limited coverage; we replace it with ATC classification. Epileptogenic/proconvulsant drug lists have an established clinical literature base [20,4,21], used to validate the hard-negative panel for epilepsy (Section 3).

Validity of bias-evaluation metrics. A smaller literature asks whether a protocol can detect what it claims to measure, not just whether a method beats a baseline: scaffold-split benchmarks overestimate performance relative to a realistic split [10], and affinity benchmarks reward training-set memorization unless reweighted [25]. Neither originates in computational biology; we know of no close RECOMB precedent. Our check sits in this family, more specific: whether a bias-diagnostic metric can see a correction's mechanism at all, a general, checkable precondition (Section 4).

3 Data

We build on the nine ChEMBL [24] indication classes used in our companion study (epilepsy, depression, antipsychotics, antibacterials, antidiabetics, antineoplastics, lipid-lowering drugs, NSAIDs, betablockers as an antihypertensive proxy), over the same 9,878-compound LINCS library, extended with nine further ATC-defined indications (Parkinson's disease, anxiolytics, antidementia drugs, opioid analgesics, renin-angiotensin antihypertensives, antacids, osteoporosis drugs, antimycotics, immunosuppressants) chosen so every original class has at least one same-area, different-mechanism negative-control partner. ATC Level-1 codes come from the ChEMBL molecule API by full InChIKey; only 954 of 9,878 compounds (9.7%) carry any ATC classification, since most of the library is research or tool compounds, not marketed drugs, why Section 5.1's enrichment is smaller than a name-matched Hub label had suggested.

SIDER [14] side-effect data is joined to the full compound set by PubChem CID, extending a prior join covering only 1,725 compounds. Coverage remains low, 543 of 9,878 (5.5%), for the same reason as ATC: SIDER only covers marketed drugs with a package insert. We use its characteristic adverse-effect terms as hard negatives (convulsion-family for epilepsy, hypotension for antihypertensives, hyperglycaemia for antidiabetics) where such a term map makes sense, not inventing one otherwise.

For epilepsy, we use a 43-drug proconvulsant list from this group's own prior, unpublished work, stated plainly rather than as an external resource: literature-curated from Chen, Albertson & Olson [4], a peer-reviewed review of drug-induced seizures, cross-validated against two public sources. Of 82 canonical compounds extracted from the original curation, 62 resolve to a compound in our library; of those, 38 (61%) carry a SIDER-listed convulsion or seizure adverse effect, and 24 (39%) match a risk tier in the independently published systematic

review of Ruffmann, Bogliun & Beghi [20]. The SIDER overlap is substantial given SIDER's own sparse 5.5% coverage, since most proconvulsants on this list are marketed drugs that resolve to SIDER at a much higher rate than the library average.

4 Method

Each compound is represented by concatenating a 32-component PCA of its ECFP4 fingerprint (TruncatedSVD, avoiding indication-matrix leakage) with a 16-dim embedding from a free logistic matrix factorization fit jointly across all 18 indications. Chemistry is necessary, not cosmetic: a pure compound-by-indication co-occurrence factorization carries almost no signal here, since 94% of the library has zero positive indications across all 18 classes and per-indication AUPRC on training data is indistinguishable from base rate. This is why NRLMF and NRBdMF add drug or disease similarity via neighborhood regularization rather than the interaction matrix alone; we add chemistry for the same reason.

On top of this shared representation, we compare four ways of scoring candidates for a target indication k_0. *Plain* fits one logistic regression per indication on the shared representation, with no area correction. *Naive subtraction* scores a candidate as the plain score minus λ times a min-max-normalized mean indication rate over k_0's negative-control set, $\lambda = 2$. *NRBdMF* factorizes a signed drug-by-disease matrix, following Azuma et al. [1]: -1 for a ChEMBL/ATC indication, $+1$ for a SIDER adverse-effect hit, chemistry withheld, matching the original specification. *Proposed* is the negative-control correction this paper is about: for k_0's negative-control set $\mathrm{NC}(k_0)$, diseases in the same ATC area treated by a different mechanism, take the train-only residuals of the already-fit plain models on $\mathrm{NC}(k_0)$'s columns, take the leading left singular vectors as an area factor A (the RUV-2 strategy of Gagnon-Bartsch and Speed [8], applied to indications rather than genes), refit k_0 with [chemistry+MF, A], and score with A's coefficients dropped.

We later tested a shrunk version of this correction, scaling A by $s \in [0, 1]$ following Gerard and Stephens' unified RUV family [9]; Section 5.5 reports no improvement.

4.1 The alignment check

Any correction scoring a candidate as a base score minus (or plus) a penalty term, naive subtraction's λ-weighted area rate and the proposed method's projected-out A-coefficients both have this form, can be probed independently of whatever metric judges it. Given a bias-diagnostic metric defined through a binary flag F (same-area membership, hard-negative membership, or any other criterion the metric penalizes), the check scores the correction's own penalty term P as a standalone classifier of F via $\mathrm{AUROC}(P, F)$. Near-0.5 AUROC means P carries no information about F: the metric cannot reward or penalize that mechanism

regardless of correction quality, so a comparison against a baseline whose term *is* informative about F is not on equal footing. We apply this to our two corrections in Section 5.7; the check names no drug-repurposing-specific quantity.

Section 5.7's result has a structural reason, stated as a proposition about how the proposed method's penalty term is built. **Proposition.** Let X be the shared representation (chemistry PCA concatenated with the free matrix-factorization embedding) on which every negative-control indication's model is fit by L_2-penalized logistic regression at inverse-strength C (class-balanced weights W), with coefficients $\hat{\beta}_j$ and fitted probabilities $\hat{P}$ on the negative-control columns, train rows only. Each residual $r_j = y_j - \hat{p}_j$ satisfies the penalized score equation $X_{\text{tr}}^{\top} W r_j = \hat{\beta}_j / C$: orthogonality to X is exact only as $C \to \infty$ (unregularized MLE); at finite C the deviation is proportional to $\hat{\beta}_j$, not zero. A, the leading left singular vectors of $R = [r_1, \ldots, r_{|\text{NC}|}]$, inherits whatever orthogonality R has.

We measured this rather than assuming it away: at the $C = 1$ used throughout, A's projection onto X's column space holds 34–40% of A's variance (`code/16_orthogonality_check.py`), not the near-zero a pure guarantee would give; the genomics domain's unregularized OLS fit (Section 5.7) is where the exact case actually holds. Whether F relates to A still depends on whether F is predictable from X: Section 5.7 tests both halves empirically.

5 Experiments

5.1 Reproducing the motivating confound

We retrained the plain ECFP4 model for epilepsy, depression, and antipsychotics, averaged over five seeds, confirming both parts of Section 1's motivating observation: bemegride ranks first of 1,522 non-antiepileptic candidates, matching the pre-existing score, and the top-50 false positives for depression and antipsychotic models are enriched for nervous-system drugs 1.5 to 1.8 times each pool's base rate (18.8% vs. 10.2%, 15.2% vs. 10.6%). We looked for, but could not locate in any saved script, the specific AUPRC collapse from 0.45 to 0.17 an earlier informal analysis attributed to naive area-score subtraction; we report what we can actually measure instead.

5.2 Evaluation protocol

We evaluate under three cross-validation schemes from the matrix-factorization repurposing literature: CVS1, a random 20% cell holdout; CVS2, a cold-start-drug holdout of 20% of compounds; CVS3, a cold-start-disease holdout of the entire target column. Each scheme runs five times with independent splits and initializations on an identical 1,702-compound pool shared by all four methods. For CVS2 we also report a scaffold-disjoint variant, since a random compound holdout can still leak near-duplicate scaffolds between train and test.

Table 1. AUPRC on true indications, CVS2 (mean over 5 reruns; 95% bootstrap CI in parentheses).

Target	Plain	Naive subtract	NRBdMF	Proposed
depression	0.201 (.114,.289)	0.190 (.105,.262)	0.039	0.181 (.112,.247)
antipsychotic	0.440 (.252,.628)	0.412 (.215,.608)	0.039	0.437 (.264,.611)
epilepsy	0.091 (.065,.127)	0.090 (.067,.128)	0.024	0.085 (.065,.112)

Table 2. Bias-diagnostic metrics. Top-50 same-area rate: lower is better (base rate ≈ 0.09–0.11). Hard-negative/proconvulsant rank: higher percentile is better.

Metric / Target	CVS	Plain	Naive subtract	NRBdMF	Proposed
area rate, depression	2	0.112	0.024	0.084	0.112
area rate, antipsychotic	2	0.136	0.096	0.084	0.144
area rate, epilepsy	2	0.144	0.068	0.120	0.176
hard-neg. pct., epilepsy	2	38.4	50.9	47.8	42.0
proconv. pct., epilepsy	2	30.9	71.8	49.2	26.3

5.3 Main results

Table 1 reports CVS2 AUPRC on the true indication label (CVS1 and CVS3 tables, same pattern, are in the repository); Table 2 reports the two bias-diagnostic metrics that matter most here, the same-ATC-area rate among the top 50 false positives and the median rank percentile of SIDER-flagged hard negatives.

Three things stand out. NRBdMF performs close to chance everywhere (AUPRC 0.02 to 0.05), confirming what the matrix diagnostics already suggested: a pure within-drug signed factorization, no chemistry, has little to work with on an 18-indication matrix this sparse. The proposed method tracks plain closely on AUPRC, the precision-preservation property it was designed for, but on both bias metrics in Table 2 it is dominated by naive subtraction, often by a wide margin; the clearest case is the proconvulsant panel, where naive pushes it from the 31st percentile to the 72nd while proposed moves it under 5 points, and in the epilepsy CVS2 row in the wrong direction. Contrary to this paper's own motivating assumption, naive subtraction's AUPRC cost is modest across every row measured under all three CVS schemes, never approaching the 62% relative collapse an earlier informal analysis had reported.

5.4 Negative-control validity check

A method that identifies an area factor from negative controls should, if those controls are mechanistically independent, leave each control's predictability roughly unchanged when corrected using a factor from the *other* controls. We ran this leave-one-negative-control-out diagnostic for every (target, control) pair; 13 of 14 cells cluster within ±0.03 of zero AUPRC change, consistent with independence. The exception, anxiolytic under antipsychotic and epilepsy targets, showed a large positive delta at 4-5 reruns (+0.127, +0.126) but also the largest std by far; re-running with 20 seeds shrank the mean to +0.034 with std essentially

flat, the signature of small-sample noise from anxiolytic's 7-positive class, not a calibration failure (full plot in the repository).

5.5 Shrinkage follow-up

The most natural explanation for Table 2 is that the negative-control classes are small, 4 to 21 positives each, and the resulting area factor is mostly noise. We tested this with empirical-Bayes-style shrinkage of the factor: $A \leftarrow sA$ for $s \in \{0, 0.25, 0.5, 0.75, 1.0\}$, selecting s and the factor rank via the same leave-one-negative-control-out diagnostic (excluding $s = 0$, which trivially achieves perfect calibration by reducing to plain). The diagnostic selected full-strength, unshrunk correction ($s = 1$) for all three targets, and re-evaluation at that setting was statistically indistinguishable from the original, ruling out noisy estimation. The more likely account is structural: the representation already absorbs much of the organ-system signal a negative-control factor would otherwise capture, or residual-SVD does not transfer well to a matrix this sparse. An RUV-4-style joint alternating estimator, re-estimating the factor and controls' fits together over several rounds, did not beat naive subtraction either; its own penalty term shows the same signature (0.496–0.502 AUROC against every flag, `results/ruv4_alignment_check.csv`), so the account holds under a mechanistically distinct strategy too.

5.6 Testing the identifiability-violation hypothesis directly

RUV's own identifying assumption [9] is that a negative-control column must carry zero loading on the covariate of interest. If that fails here, the area factor would absorb some of the target's own signal, and projecting it out would cost real precision along with the confound. For each target and CVS scheme where the area factor is fit, we regressed the target's true label, train rows only, onto its fitted area factor, and measured R^2, which should sit near zero under the assumption. Across all six cells the values are small, 0.0006 to 0.021, and do not track naive's advantage: epilepsy has the smallest overlap but the largest gap in naive's favor, the opposite of what the hypothesis predicts. This rules it out.

5.7 Applying the alignment check

Ruling out identifiability leaves a more mundane candidate: the bias metrics may sit closer to what naive subtraction directly penalizes, raw negative-control membership, than to what the proposed method targets, a residualized factor estimated after chemistry and matrix-factorization signal are removed. We apply the alignment check of Section 4.1: for each target, under CVS2, we compute naive's penalty term (raw negative-control rate) and proposed's penalty term made concrete (the contribution A's coefficients would have made had they not been projected out) for every non-indicated test compound, and score each, on its own, as a classifier of the exact compounds the bias metrics flag: same-ATC-area membership, SIDER-hard-negative, or proconvulsant-panel membership.

Table 3. AUROC of each method's penalty term, and of X alone, as a classifier of the flagged compounds (mean ± std, 5 seeds, CVS2). 0.5 = no information.

Target	Flag	Naive's term	Proposed's term	X alone
depression	same-ATC-area	0.766 ± 0.051	0.499 ± 0.002	0.718 ± 0.015
depression	hard-negative	0.636 ± 0.025	0.499 ± 0.002	0.514 ± 0.066
antipsychotic	same-ATC-area	0.774 ± 0.046	0.498 ± 0.005	0.706 ± 0.038
antipsychotic	hard-negative	0.663 ± 0.033	0.501 ± 0.002	0.591 ± 0.065
epilepsy	same-ATC-area	0.787 ± 0.055	0.497 ± 0.007	0.722 ± 0.040
epilepsy	hard-negative	0.600 ± 0.042	0.500 ± 0.002	0.613 ± 0.040
epilepsy	proconvulsant	0.721 ± 0.084	0.500 ± 0.002	0.695 ± 0.081

Table 3 is the resolution, not a further hint. Naive's penalty term separates every bias-flagged compound set well above chance, 0.60 to 0.79 AUROC, across every target. The proposed method's penalty term sits at 0.497 to 0.501 in every row, with a standard deviation of 0.002 to 0.007, too tight for noise to explain the gap, though cleanest below for same-ATC-area and weaker for hard-negative/proconvulsant, where X predicts the flag less strongly. By the Proposition in Section 4.1, the proposed method's penalty term is only partially orthogonal, on train rows, to the chemistry-plus-matrix-factorization representation X that every per-indication model is fit on, measured rather than assumed; near-chance AUROC against flag F still tracks whether F is predictable from X itself. The rightmost column tests that premise rather than assuming it, by fitting a plain classifier on X alone, train rows only, to predict each flag on the same candidate population. The same-ATC-area flag is well predicted by X on every target, 0.706 to 0.722, close to naive subtraction's own 0.766 to 0.787: most of what makes a compound look like it belongs to the confounded area is already present in chemistry and the shared matrix-factorization embedding before any negative-control information is added. The hard-negative and proconvulsant flags are predictable from X too, but less uniformly, from 0.514 for depression's hard negatives (indistinguishable from chance) up to 0.695 for the epilepsy proconvulsant panel. This account rests on measurement, not a pure guarantee, in this domain; Section 5.7's genomics domain is where the guarantee is exact.

This two-step recipe, orthogonality by construction, then an empirical check of whether the metric's flag is predictable from the correction's base representation, names nothing specific to drug repurposing, ATC codes, or RUV, and applies wherever a residualized correction meets a bias metric built on group membership the same base features already predict.

We also re-ran this check outside drug repurposing, on the dataset RUV-2 was introduced on: Vawter et al.'s brain gene-expression study [22], gender as the factor of interest, lab as the confound, and Gagnon-Bartsch and Speed's 799 housekeeping controls [7]. Step one holds exactly: each control gene's OLS residual on gender satisfies the normal equations to floating-point precision, so A's genomics analog is exactly orthogonal to gender. Step two differs: lab is not predictable from gender (AUROC 0.39–0.41, near or below chance), unlike same-ATC-area membership above. As predicted, the orthogonalized factor stays

informative about lab (0.50–0.94 alone, 0.95–1.00 combined), matching or beating naive's unprojected summary (0.55–0.76): the two domains diverge exactly where the check predicts. Donor-grouped CV leaves these numbers unchanged.

A third, harder domain: bladderbatch [6,15], a 3-level cancer-status factor fit by ridge regression against a 5-level batch confound (57 samples, one level $n = 4$, a thin base). Step one again holds exactly by the normal equation. Step two differs: cancer predicts batch moderately (0.722), yet the orthogonalized factor stays as informative about batch as naive's raw mean, or more so (0.68 against 0.62), not falling to chance, likely because X's 2 columns are too low-dimensional to absorb batch's variance at moderate correlation, unlike the 48-dimensional drug-repurposing case. This refines, rather than confirms, the check's second step: predictability of F from X matters less than whether X is rich enough to exhaust it.

6 Discussion and Limitations

This paper set out to build a more principled alternative to a correction assumed too costly to use, found instead that naive wins on every bias metric measured, and asked whether that scoreboard meant what it appeared to. It did not: the Proposition in Section 4.1 shows the correction's penalty term is orthogonal to its base representation by construction only without regularization, and only partial orthogonality survives here, measured rather than assumed; Table 3's rightmost column shows the flags are substantially predictable from that same representation regardless. A term built to contain little of what the representation explains has correspondingly little left to say about a flag it already explains well. This is not post-hoc: Sections 5.6 and 5.7 ruled out an identifiability violation before confirming the alignment account.

We think the check itself is the more durable contribution, with one caveat: orthogonality of a residual to its design matrix is standard GLM theory, not new; whether a flag is predictable from that matrix is empirical, and we got this wrong ourselves at first, stating orthogonality as unconditional before measuring the real 34–40% violation (Section 4). Negative-control identification of a latent confound is established in genomics [8,9,16] and pharmacoepidemiology [18,23]; comparing it against a simpler baseline can mislead when the flag aligns more with the baseline's term, which Section 5.7's recipe catches first. NRBdMF [1] runs chemistry-withheld, matching its original specification, so its near-floor AUPRC reflects weak co-occurrence signal alone (Section 3), not a fair baseline; it also has no residualized penalty term, so the check does not apply.

Several limitations remain. The check spans three domains (Section 5.7), the third partial; its empirical half must be re-run per case. ATC/SIDER cover under 10% of the library. CVS3 is degenerate for target-specific methods, collapsing plain and proposed together. Table 3 measures each term against the metrics' flags, not ground truth; we skipped repoDB [2] and did not fully read DRDM [19]. Code: `https://github.com/edanmn/area-debiasing-drug-repurposing`.

References

1. Azuma, I., Mizuno, T., Kusuhara, H.: NRBdMF: A recommendation algorithm for predicting drug effects considering directionality. Journal of Chemical Information and Modeling **63**(2), 474–483 (2023). https://doi.org/10.1021/acs.jcim.2c01210
2. Brown, A.S., Patel, C.J.: A standard database for drug repositioning. Scientific Data **4**, 170029 (2017). https://doi.org/10.1038/sdata.2017.29
3. Campillos, M., Kuhn, M., Gavin, A.C., Jensen, L.J., Bork, P.: Drug target identification using side-effect similarity. Science **321**(5886), 263–266 (2008). https://doi.org/10.1126/science.1158140
4. Chen, H.Y., Albertson, T.E., Olson, K.R.: Treatment of drug-induced seizures. British Journal of Clinical Pharmacology **81**(3), 412–419 (2016)
5. Corsello, S.M., Bittker, J.A., Liu, Z., Gould, J., McCarren, P., Hirschman, J.E., Johnston, S.E., Vrcic, A., Wong, B., Khan, M., Asiedu, J., Narayan, R., Mader, C.C., Subramanian, A., Golub, T.R.: The drug repurposing hub: A next-generation drug library and information resource. Nature Medicine **23**(4), 405–408 (2017). https://doi.org/10.1038/nm.4306
6. Dyrskjot, L., Kruhøffer, M., Thykjaer, T., Marcussen, N., Jensen, J.L., Møller, K., Ørntoft, T.F.: Gene expression in the urinary bladder. Cancer Research **64**(11), 4040–4048 (2004)
7. Eisenberg, E., Levanon, E.Y.: Human housekeeping genes are compact. Trends in Genetics **19**(7), 362–365 (2003)
8. Gagnon-Bartsch, J.A., Speed, T.P.: Using control genes to correct for unwanted variation in microarray data. Biostatistics **13**(3), 539–552 (2012). https://doi.org/10.1093/biostatistics/kxr034
9. Gerard, D., Stephens, M.: Unifying and generalizing methods for removing unwanted variation based on negative controls. Statistica Sinica **31**, 1145–1166 (2021)
10. Guo, Q., Hernandez-Hernandez, S., Ballester, P.J.: Scaffold splits overestimate virtual screening performance (2024)
11. Johnson, W.E., Li, C., Rabinovic, A.: Adjusting batch effects in microarray expression data using empirical bayes methods. Biostatistics **8**(1), 118–127 (2007). https://doi.org/10.1093/biostatistics/kxj037
12. Kondadadi, R.K., McCreary, D.: How much do evaluation choices inflate transcriptomic drug repurposing? a measured audit on LINCS L1000 (2026), manuscript in preparation
13. Kuhn, M., Campillos, M., Letunic, I., Jensen, L.J., Bork, P.: A side effect resource to capture phenotypic effects of drugs. Molecular Systems Biology **6**, 343 (2010). https://doi.org/10.1038/msb.2009.98
14. Kuhn, M., Letunic, I., Jensen, L.J., Bork, P.: The SIDER database of drugs and side effects. Nucleic Acids Research **44**(D1), D1075–D1079 (2016). https://doi.org/10.1093/nar/gkv1075
15. Leek, J.T., Scharpf, R.B., Corrada Bravo, H., Simcha, D., Langmead, B., Johnson, W.E., Geman, D., Baggerly, K., Irizarry, R.A.: Tackling the widespread and critical impact of batch effects in high-throughput data. Nature Reviews Genetics **11**, 733–739 (2010)
16. Leek, J.T., Storey, J.D.: Capturing heterogeneity in gene expression studies by surrogate variable analysis. PLoS Genetics **3**(9), e161 (2007). https://doi.org/10.1371/journal.pgen.0030161

17. Lin, G., Zhang, X., Ren, Z., Zou, Q., Tiwari, P., Zhou, C., Ding, Y.: TAPB: An interventional debiasing framework for alleviating target prior bias in drug-target interaction prediction. Nature Communications **16** (2025). https://doi.org/10.1038/s41467-025-66915-1
18. Lipsitch, M., Tchetgen Tchetgen, E., Cohen, T.: Negative controls: A tool for detecting confounding and bias in observational studies. Epidemiology **21**(3), 383–388 (2010). https://doi.org/10.1097/EDE.0b013e3181d61eeb
19. Meng, Y., Wang, Y., Hu, X., Lu, C., Tang, X., Cui, F., Zeng, P., Yao, Y., Yang, J., Xu, J.: Adaptive debiasing learning for drug repositioning. Journal of Biomedical Informatics **167**, 104843 (2025). https://doi.org/10.1016/j.jbi.2025.104843
20. Ruffmann, C., Bogliun, G., Beghi, E.: Epileptogenic drugs: A systematic review. Expert Review of Neurotherapeutics **6**(4), 575–589 (2006)
21. Shorvon, S., Tallis, R., Wallace, H.: Antiepileptic drugs: Coprescription of proconvulsant drugs and oral contraceptives: A national study of antiepileptic drug prescribing practice. Journal of Neurology, Neurosurgery & Psychiatry **72**, 114–115 (2002)
22. Vawter, M.P., Evans, S., Choudary, P., Tomita, H., Meador-Woodruff, J., Molnar, M., Li, J., Lopez, J.F., Myers, R., Cox, D., Watson, S.J., Akil, H., Jones, E.G., Bunney, W.E.: Gender-specific gene expression in post-mortem human brain: Localization to sex chromosomes. Neuropsychopharmacology **29**(2), 373–384 (2004)
23. Zafari, Z., Park, J.e., Shah, C.H., dosReis, S., Gorman, E.F., Hua, W., Ma, Y., Tian, F.: The state of use and utility of negative controls in pharmacoepidemiologic studies. American Journal of Epidemiology **193**(3), 426–453 (2023). https://doi.org/10.1093/aje/kwad201
24. Zdrazil, B., Felix, E., Hunter, F., Manners, E.J., Blackshaw, J., Corbett, S., de Veij, M., Ioannidis, H., Lopez, D.M., Mosquera, J.F., Magarinos, M.P., Bosc, N., Arcila, R., Kizilören, T., Gaulton, A., Bento, A.P., Adasme, M.F., Monecke, P., Landrum, G.A., Leach, A.R.: The ChEMBL database in 2023: A drug discovery platform spanning multiple bioactivity data types and time periods. Nucleic Acids Research **52**(D1), D1180–D1192 (2024). https://doi.org/10.1093/nar/gkad1004
25. Zhang, C., Hu, Z., Jiang, C., Chen, W., Xu, J., Zhang, S.: Rethinking the generalization of drug-target affinity prediction algorithms via similarity-aware evaluation. In: International Conference on Learning Representations (ICLR) (2025)